

1

point

1.
- Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors should be 10000 dimensional, so as to capture the full range of variation and meaning in those words.
- ☐

True
- ☐

False

1

point

2.
- What is t-SNE?
- ☐

A linear transformation that allows us to solve analogies on word vectors
- ☐

A non-linear dimensionality reduction technique
- ☐

A supervised learning algorithm for learning word embeddings
- ☐

An open-source sequence modeling library

1

point

3.
- Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word “ecstatic” does not appear in your small training set, your RNN might reasonably be expected to recognize “I’m ecstatic” as deserving a label $y = 1$.

- ☐

True
- ☐

False

1

point

4.
- Which of these equations do you think should hold for a good word embedding? (Check all that apply)
- ☐

 $e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$
- ☐

 $e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$
- ☐

 $e_{boy} - e_{brother} \approx e_{girl} - e_{sister}$
- ☐

 $e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$

1

point

5.
- Let E be an embedding matrix, and let o_{1234} be a one-hot vector corresponding to word 1234. Then to get the embedding of word 1234, why don't we call $E * o_{1234}$ in Python?
- ☐

It is computationally wasteful.
- ☐

The correct formula is $E^T * o_{1234}$.
- ☐

This doesn't handle unknown words (<UNK>).
- ☐

None of the above: calling the Python snippet as described above is fine.

1

point

6.
- When learning word embeddings, we create an artificial task of estimating $P(target \mid context)$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.
- ☐

True
- ☐

False

1

point

7.
- In the word2vec algorithm, you estimate $P(t \mid c)$, where t is the target word and c is a context word. How are t and c chosen from the training set? Pick the best answer.
- ☐

 c is a sequence of several words immediately before t .
- ☐

 c and t are chosen to be nearby words.
- ☐

 c is the one word that comes immediately before t .
- ☐

 c is the sequence of all the words in the sentence before t .

1

point

8.
- Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{i=1}^{10000} e^{\theta_i^T e_c}}$$

Which of these statements are correct? Check all that apply.

- ☐

 θ_t and e_c are both 500 dimensional vectors.
- ☐

 θ_t and e_c are both 10000 dimensional vectors.
- ☐

 θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.
- ☐

After training, we should expect θ_t to be very close to e_c when t and c are the same word.

1

point

9.
- Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings.The GloVe model minimizes this objective:

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - \log X_{ij})^2$$

Which of these statements are correct? Check all that apply.

- ☐

 θ_i and e_j should be initialized to 0 at the beginning of training.
- ☐

 θ_i and e_j should be initialized randomly at the beginning of training.
- ☐

 X_{ij} is the number of times word i appears in the context of word j.
- ☐

The weighting function $f(\cdot)$ must satisfy $f(0) = 0$.

1

point

10.
- You have trained word embeddings using a text dataset of m_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of m_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstance would you expect the word embeddings to be helpful?
- ☐

 $m_1 \gg m_2$
- ☐

 $m_1 \ll m_2$

Upgrade to submit